



Databases & Schema Design

Intro to MongoDB, Object Database Mappers and Schema Design

Outline

- History of Database
- Why MongoDB
- Mongoose — Connecting MongoDB and Node.js
- Tooling
- Workshop

What is a Database?

- Excel
- Cellphone directory
- Inbox
- Folder System
- The Web

Database Management Sys. (DBMS)

- How to define records (Data Definition)
- Creating/Updating/Deleting Records
- Querying/Retrieval of Records
- Administering the database system

Progression of Databases

- Navigational (< 1970s)
- Relational (> 1970s)
- NoSQL (> 2000s)

MongoDB History

- Founded in 2007 as 10gen
- Engine for Google Apps
- Open Source, but driven by the company MongoDB
- Puts the M in MEAN

NoSQL: MongoDB

- ◉ Not Only SQL — Can have SQL-like queries, but:
- ◉ Does not restrict to storing info in tables (relations)
- ◉ Stores information in BSON (Binary JSON) *documents*
 - JSON (JavaScript Object Notation) is a subset of actual JS objects:
 - ```
{
 "firstName" : "Frodo",
 "age" : 33,
 "possessions" : ["ring", "sword"]
}
```
  - Must use double quotes, including around key names.
  - BSON is a superset of JSON which is stored as binary values, including some additional data types (e.g. integers, dates) and metadata (number of bits)

# Collections vs. Tables

- The Mongo equivalent of a table is a *collection*
- Collections don't have schemas — docs can vary



# Documents vs. Rows

- Each piece of data in Mongo is called a Document
- Not PDF or DOC files, but like JS objects or associative arrays
- A Document is a JSON object, like JavaScript objects
- Each Document can have its own design



# Mongo vs. SQL vs. Excel

| Excel      | SQL             | Mongo            |
|------------|-----------------|------------------|
| Excel File | Database        | Database         |
| Worksheet  | Table           | Collection       |
| Row        | Row             | Document         |
| Column     | Column          | Field            |
|            | Index           | Index            |
| VLOOKUP    | Join            | Embedding        |
|            | Primary Key: id | Primary Key: _id |

# SQL to NoSQL

| SQL Term    | Mongo NoSQL Term    |
|-------------|---------------------|
| Database    | Database            |
| Table       | Collection          |
| Row         | Document            |
| Column      | Field               |
| Index       | Index               |
| Join        | Embedding & Linking |
| Primary Key | _id field           |
| Group By    | Aggregation         |



# Talking to Mongo using JS

- Talk to Mongo using JavaScript functions
- The “db” object is the pathway to the DB

| SQL                                                     | Mongo Statement                               |
|---------------------------------------------------------|-----------------------------------------------|
| CREATE TABLE users (name varchar, age int)              | Done automatically while inserting            |
| INSERT INTO users (name, age) VALUES ("Nicole", 24)     | db.users.insert({"name":"Nicole", "age":24}); |
| SELECT * FROM users;                                    | db.users.find()                               |
| SELECT * FROM users WHERE age=24;                       | db.users.find({"age":24})                     |
| SELECT * FROM users WHERE name="Jay" ORDER BY name ASC; | db.users.find({"name":"Jay"}).sort("name":1)  |

# MongoDB Strengths

- ◉ **Schema-less (no enforced document structure)**
  - Collections can have mixed docs!
  - Agile development - changing, flexible
  - Capturing unstructured information
- ◉ **Preplanning for scale**
- ◉ **Use cases:**
  - Archiving / event logging (Hummingbird)
  - Document / content management (Forbes)
  - Gaming (EA FIFA)
  - Mobile & location-based (Foursquare)
  - Realtime statistics (Criticism)

# MongoDB Features

- **Geospatial data type — Foursquare**
  - Spherical searches
  - Inclusion search (within polygons)
- **Horizontal scaling (replica sets, sharding)**
- **Aggregation (map-reduce)**
- **Stored JS functions**
- **Low JS impedance**
- **Rich query capabilities**

# MongoDB Weaknesses

- MongoDB is consistent at the *document* level, but does not have ACID-compliant multi-document *transactions*
- ACID (key in banking / accounting)
  - Atomic
  - Consistent
  - Isolated
  - Durable
- Transactions
  - Cannot change multiple docs at the same time in an all-or-none fashion
- Maybe you want SQL

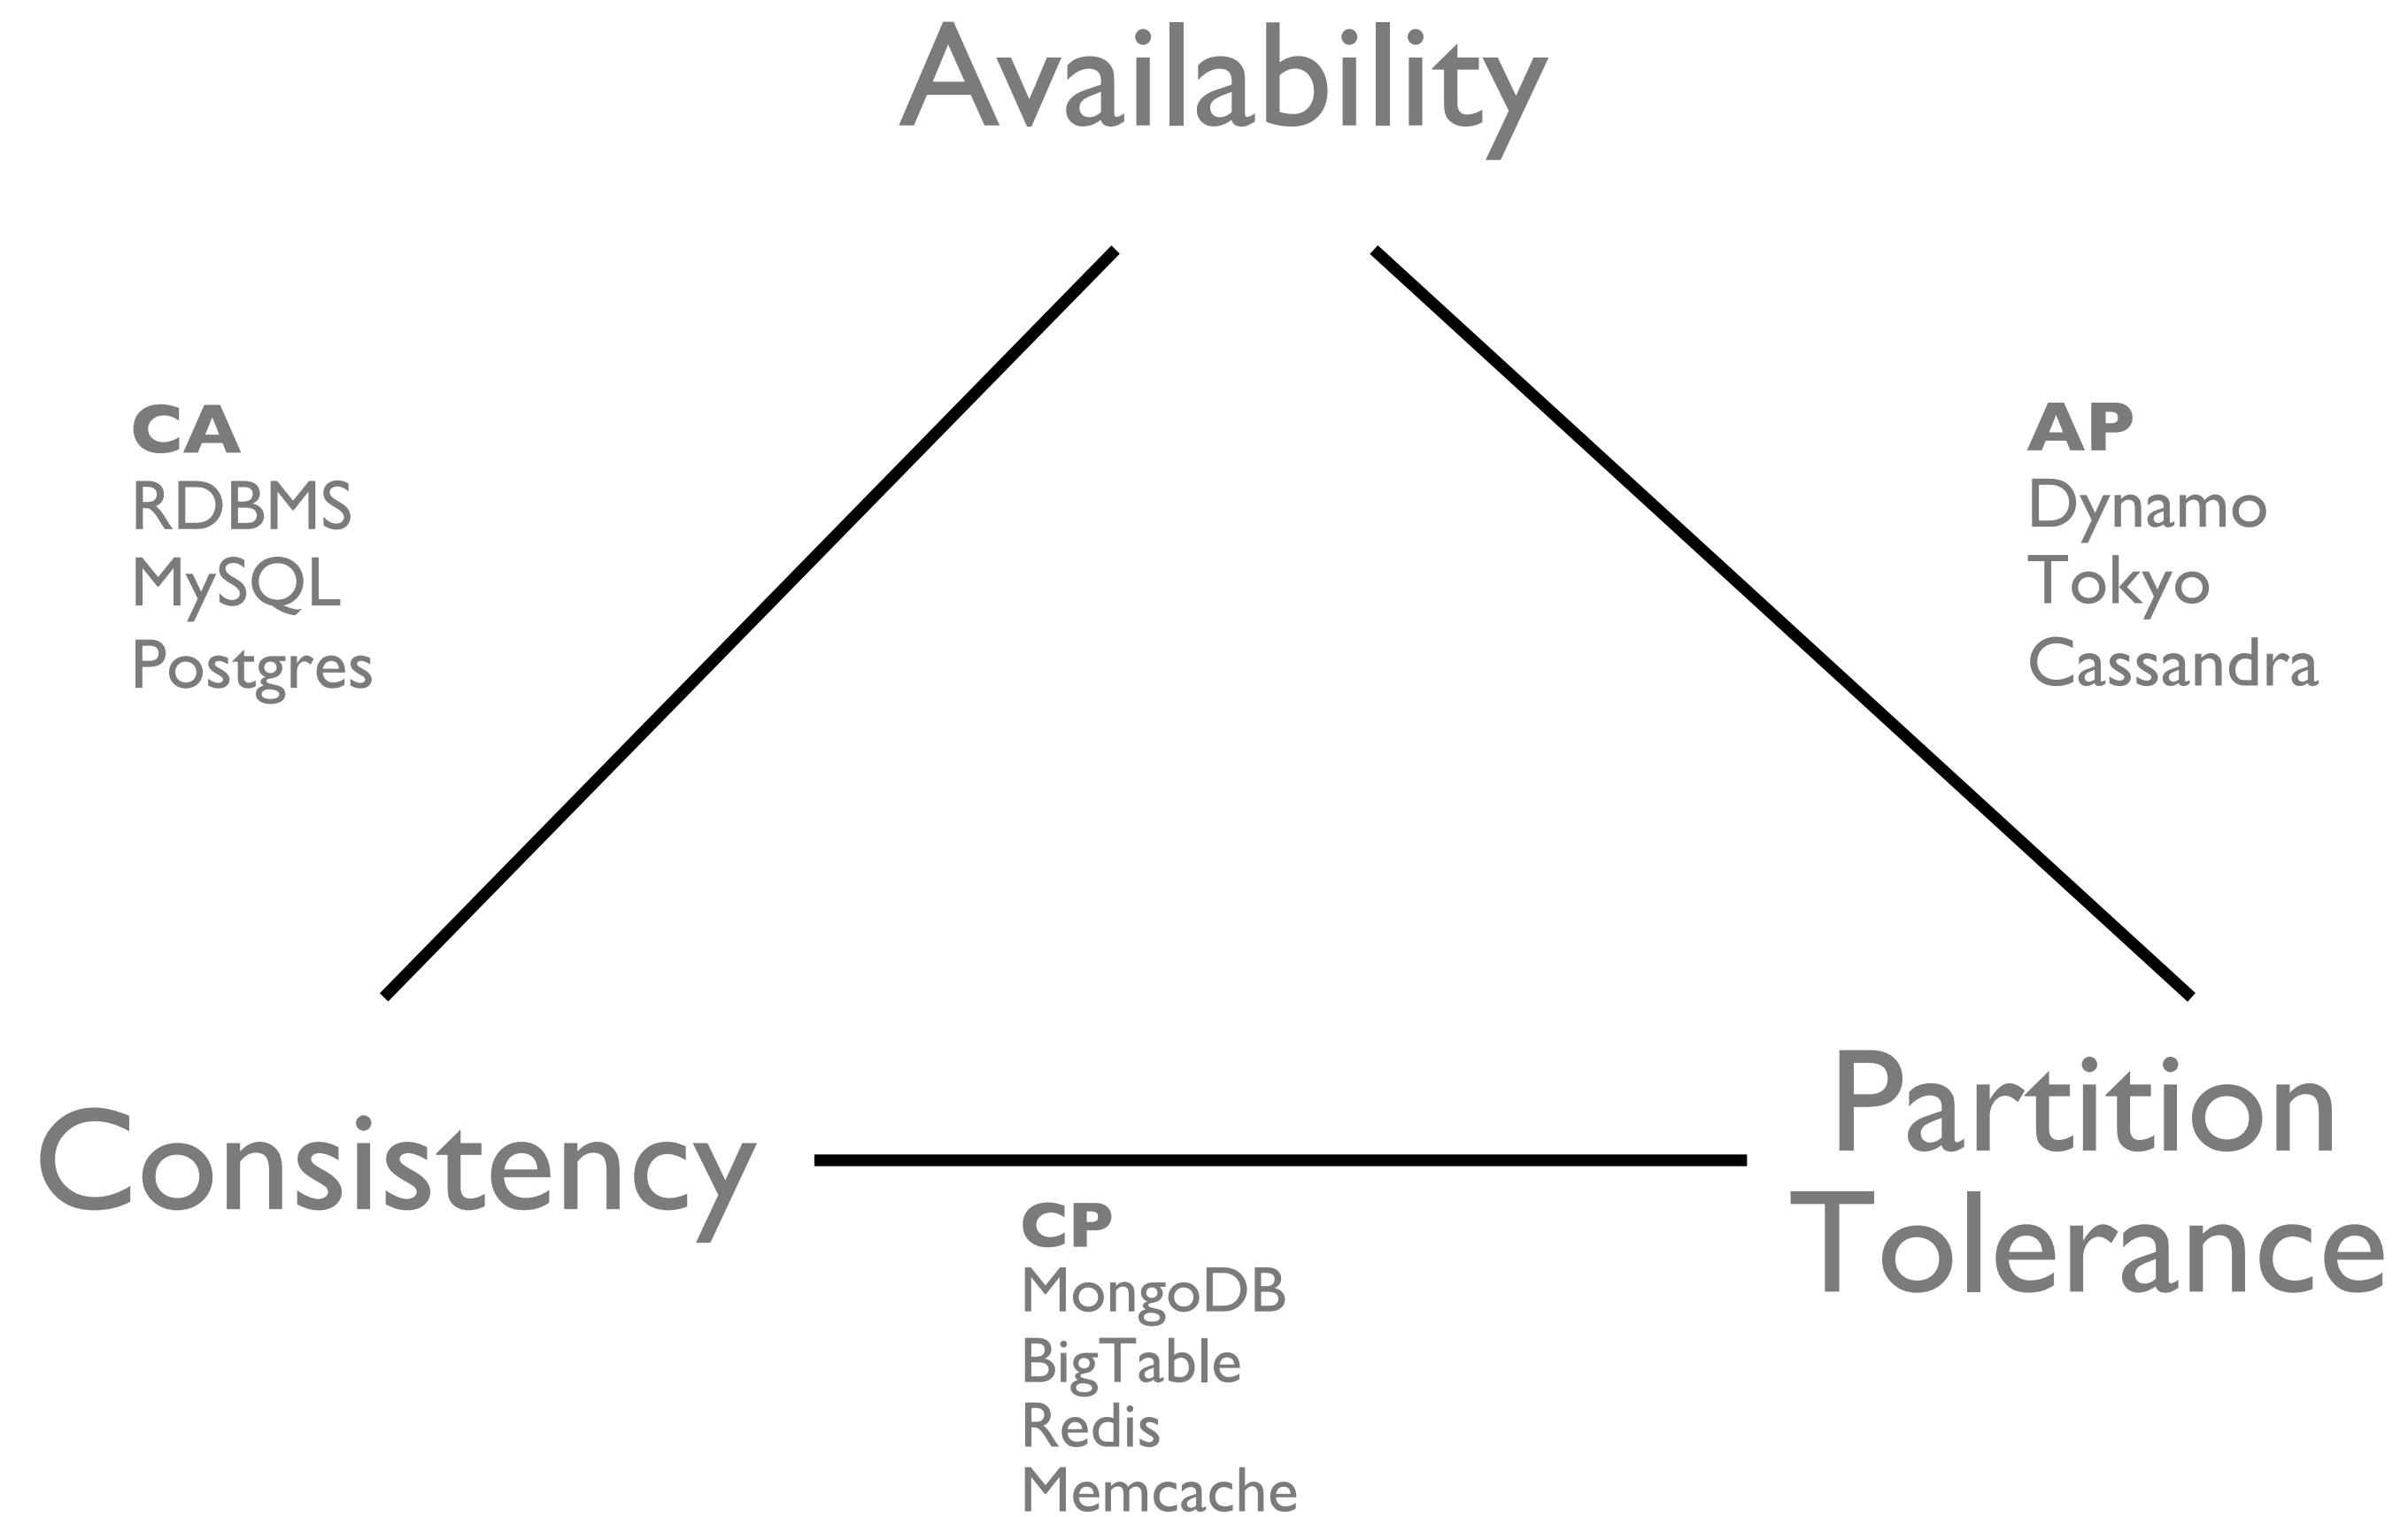
# CAP "Theorem"

- **Consistency:** all clients see the same data
- **Availability:** each client can always read or write
- **Partition Tolerance:** system works even if network splits
- **"Pick 2"**
  - (Reality is more complex...)





# CAP, metaphorically speaking



# CRUD in Mongo

- Create
- Read / Query
- Update
- Delete



# Create/Insert

```
use demodatabase;
db.users.insert({
 firstName: "David",
 lastName: "Yang",
 username: "davidyang",
 password: "74824h234b280412312fdgsdgae"
});
```





# Read

```
> db.users.find()
[{
 "_id" : ObjectId("52f1014ace52721056688d7a"),
 "firstName" : "David",
 "lastName" : "Yang",
 "username" : "davidyang",
 "password" : "74824h234b280412312fdgsdgae"
}]

> db.users.findOne()
{
 "_id" : ObjectId("52f1014ace52721056688d7a"),
 "firstName" : "David",
 "lastName" : "Yang",
 "username" : "davidyang",
 "password" : "74824h234b280412312fdgsdgae"
}
```



# Update

```
db.users.update(queryObject, updateCommandObject);
```

```
db.users.update(
 { "username": "david" }, // query object
 { "username": "david", "favorite_movies" : [
 "Braveheart", "Terminator 2"
] } // update command
)
```

```
db.users.update(
 { "username": "david" },
 { $set: { "favorite_movies": ["Braveheart", "Johnny Bravo"] } }
)
```



# Delete

```
db.users.remove({ _id: ObjectId("52f1014ace52721056688d7a")})
```

```
> db.users.count()
0
```

# Querying

- **Matchers**
- **Range (\$gt, \$lt, \$gte, \$lte)**
- **Set Operators (\$in, \$nin, \$all)**
- **Boolean Expressions (\$or,**
- **Arrays**
- **Regular Expressions (/regex/)**
- **JavaScript (functions)**

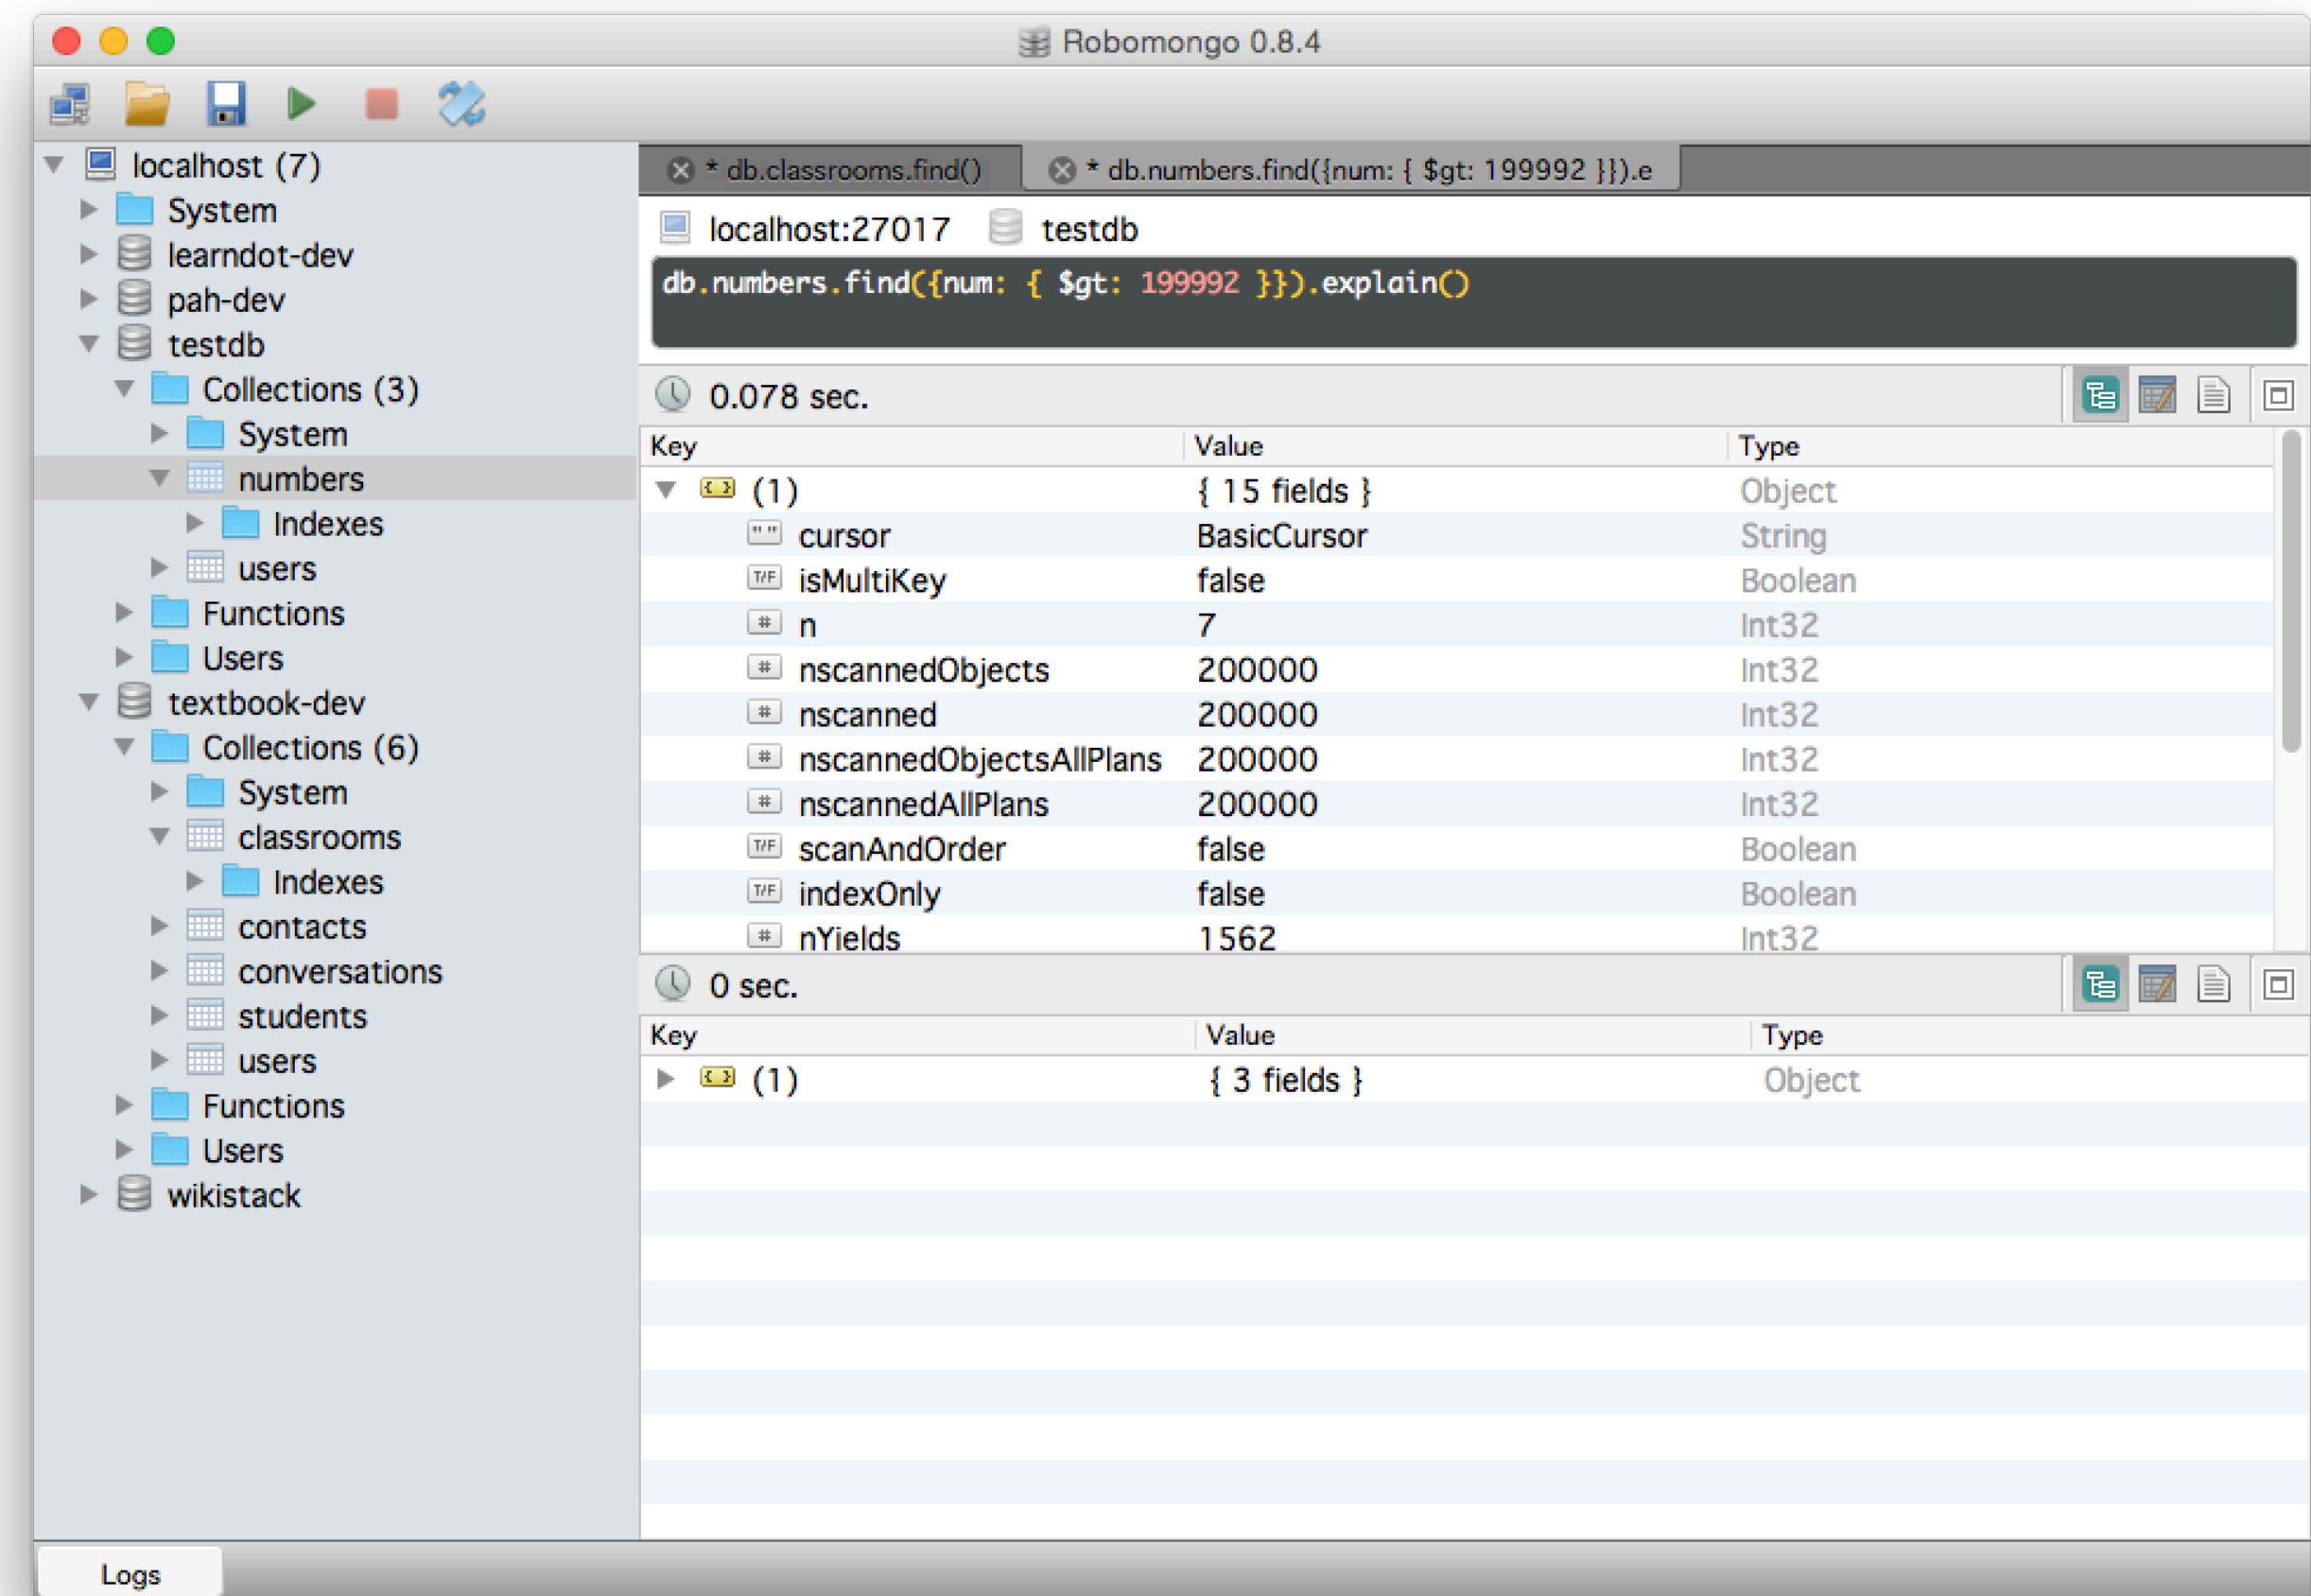


# Indexes

- **Indexes are the primary tool DB administrators have to improve performance of querying**
- **Like the index of book, without one, any time you want to find something, you have to scan the entire table**



# Tooling


## ● Robomongo







# Tooling



## ● Mongohacker

 This repository Search Explore Gist Blog Help  davidy



 TylerBrock / mongo-hacker Watch 46 ★ S









MongoDB Shell Enhancements for Hackers <http://tylerbrock.github.io/mongo-hacker>


 193 commits  2 branches  4 releases  26 contributors

 branch: master **mongo-hacker** / + 

Merge pull request #113 from tobsn/patch-1 ...

 TylerBrock authored on Jan 9 latest commit 3e776b455b 

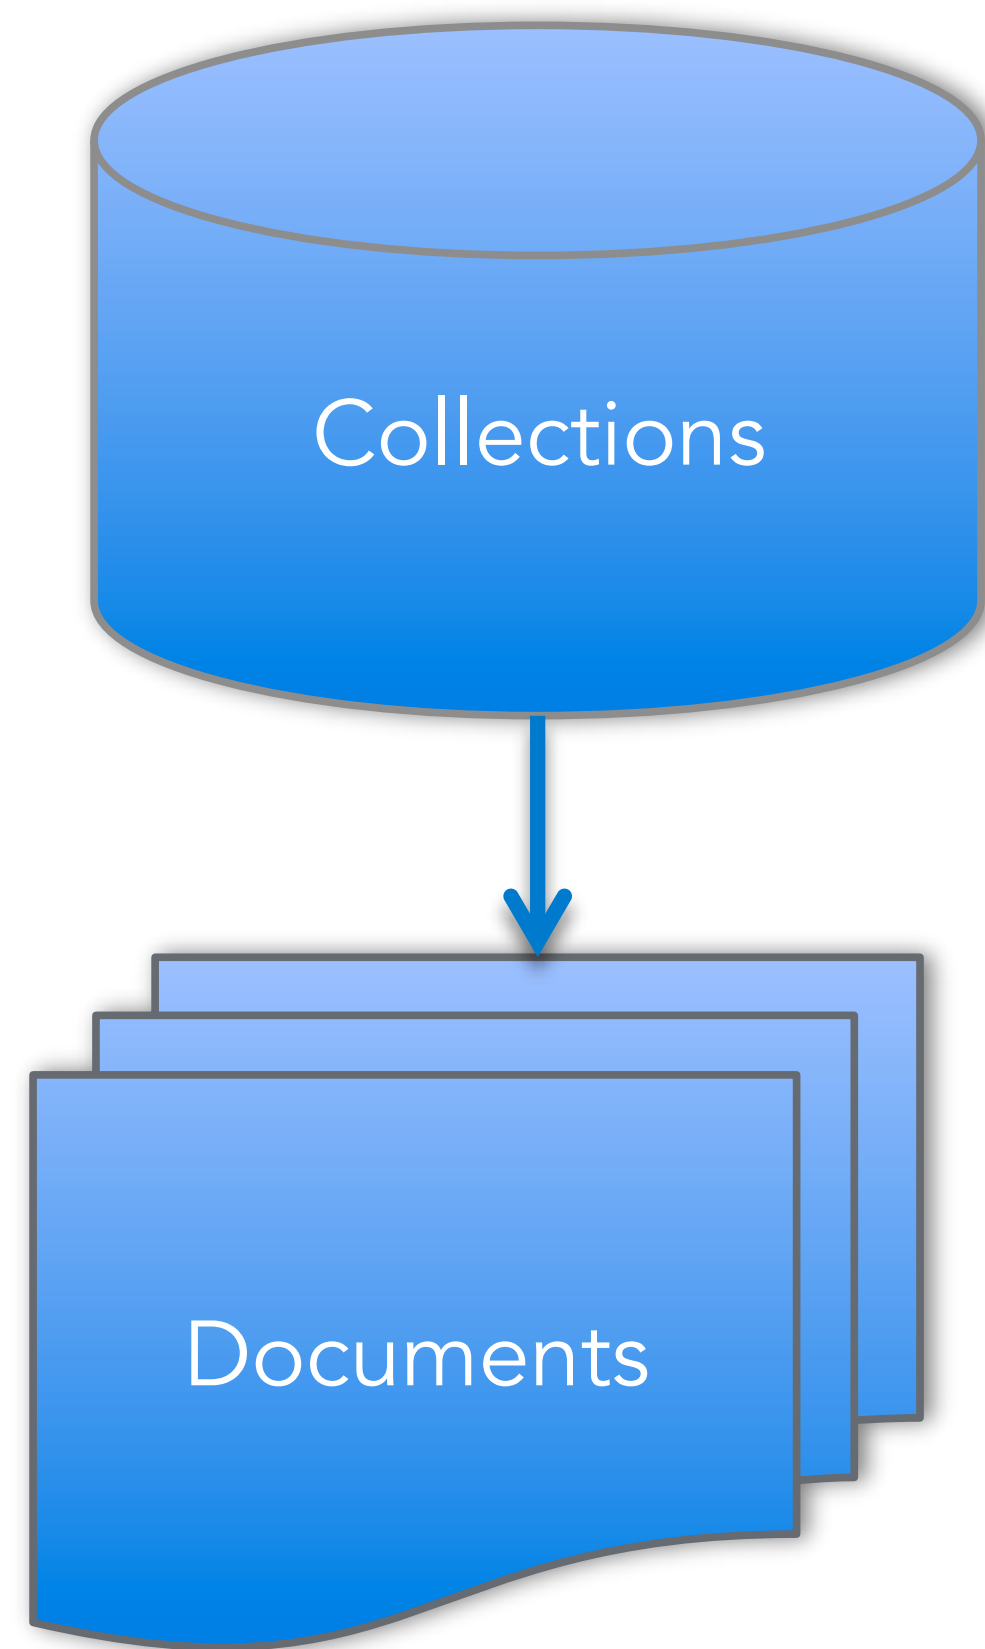
|                                                                                                    |                                                                     |              |
|----------------------------------------------------------------------------------------------------|---------------------------------------------------------------------|--------------|
|  hacks        | changed n._limit to n / n is an integer and has no _limit attribute | 2 months ago |
|  .gitignore   | adding built mongo_hacker output to gitignore                       | 2 years ago  |
|  LICENSE      | Added license file                                                  | 5 months ago |
|  Makefile     | minor: restore original .mongorc.js on uninstall                    | 4 months ago |
|  README.md    | Update README.md                                                    | 2 months ago |
|  base.js      | Fix #76: Add Support for windows                                    | 5 months ago |
|  config.js    | change default key color to grey to increase visibility             | 4 months ago |
|  package.json | bump 0.0.4                                                          | 2 months ago |

 README.md

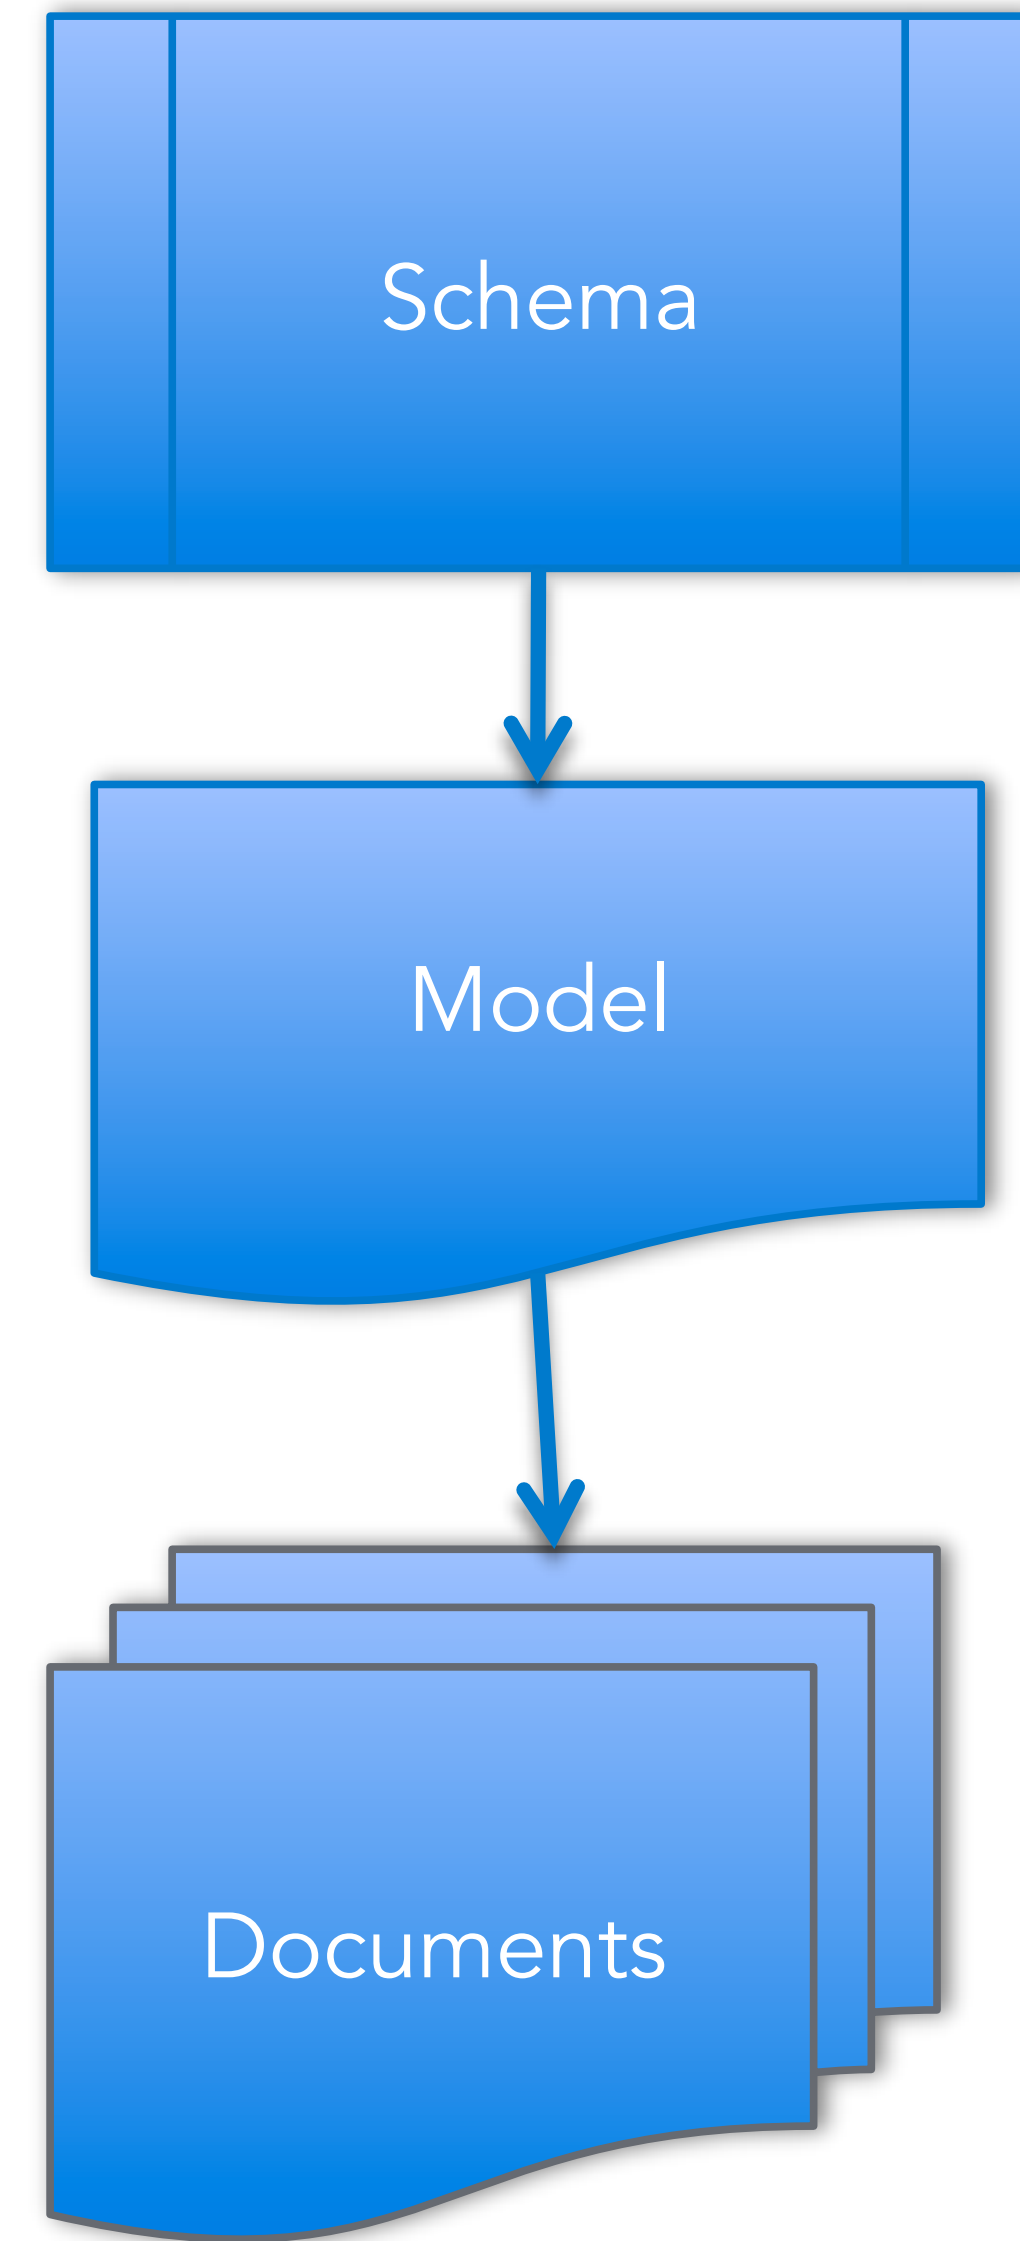
## MongoDB Shell Enhancements

# Mongoose

- **Mongoose is an Object-Document Mapper that makes it easy to access MongoDB from Node.js**
- **Mongoose Features:**
  - Schema modeling/validation
  - Data casting (convert mongo types to JS types)
  - Query building
  - Hooks (code that runs pre/post save/delete/update)
  - Class and instance methods of models
  - Getters and setters



=



# Mongoose Basics

- Make a **Schema** (interactive blueprint object)
- Extend the **Schema** with **Hooks**, **Methods**, **Statics**, **Virtuals**, etc.
- Use the completed **Schema** to build a **Model** (object that allows interacting with a Mongo collection — also, can act a constructor function)
- Use the **Model** (Collection) to create/find **Instances**
- Use the **Instances** (individual Docs) to save/update/delete



# Create a Schema

```
var mongoose = require('mongoose');
var Schema = mongoose.Schema;

var blogSchema = new Schema({
 title: String,
 author: String,
 body: String,
 comments: [{body: String, date: Date }],
 date: {type: Date, default: Date.now },
 hidden: Boolean,
 meta: {votes: Number, favs: Number }
});
```



# Build Model from Schema

```
var Blog = mongoose.model('Blog', blogSchema);
```





# Model & Instance usage

```
var myBlog = new Blog({
 title: "David's Blog",
 author: "David"
});

myBlog.save(function(err, myBlog) {
 if(err) console.log("Error saving", err);
 console.log(myBlog);
});

Blog.find({ title: /David/ });
```

# Mongoose

- Lives inside Node.js process
- knows how to communicate via the MongoDB Binary Protocol to the MongoDB Server

# Schema Design

- **How do we store data in Mongo?**
  - Non-relational
  - No support for Database Joins
- **Since MongoDB documents are just JSON documents they support array and object embedding**
  - New choice in MongoDB: do I embed this document or reference?
  - MongoDB doesn't support JOINS, you have to do it in your JavaScript application

# Background: SQL & Joins

- **Three Basic Principles in SQL**
  - Data is tabular
  - Each cell of the table contains one value (no arrays/collections/objects)
  - No redundancy of data (each row represents the single-source of truth about that data)
- **Storing multiple data points (relations)**
  - Avoid redundancy / multiple columns / complex updates using join tables
  - Cost is that performing a join is slow

# Embedding

```
var mongoose = require('mongoose');
var Schema = mongoose.Schema;
```

```
var blogSchema = new Schema({
 title: String,
 body: String,
 comments: [{
 body: String,
 date: Date
 }]
});
```

OR

```
var mongoose = require('mongoose');
var Schema = mongoose.Schema;
```

```
var commentSchema = new Schema({
 body: String,
 date: Date
});
```

```
var blogSchema = new Schema({
 title: String,
 body: String,
 comments: [commentSchema]
});
```

# Referencing

```
var mongoose = require('mongoose');
var Schema = mongoose.Schema;
var commentSchema = new Schema({
 body: String,
 date: Date
});
var Comment = mongoose.model('Comment', commentSchema);
var blogSchema = new Schema({
 title: String,
 body: String,
 comments: [{
 type: Schema.Types.ObjectId,
 ref: 'Comment'
 }]
});
```

# Embedding vs. Referencing

```
var mongoose = require('mongoose');
var Schema = mongoose.Schema;
```

```
var blogSchema = new Schema({
 title: String,
 body: String,
```

```
 comments: [{
 body: String,
 date: Date
 }]
});
```

```
comments: [{
 body: String,
 date: Date
}]
```

```
var mongoose = require('mongoose');
var Schema = mongoose.Schema;
var commentSchema = new Schema({
 body: String,
 date: Date
});
var Comment = mongoose.model('Comment', commentSchema);
var blogSchema = new Schema({
 title: String,
 body: String,
```

```
 comments: [{
 type: Schema.Types.ObjectId,
 ref: 'Comment'
 }]
});
```

```
comments: [{
 type: Schema.Types.ObjectId,
 ref: 'Comment'
}]
```

# Schema Choices

- **Cardinality**
  - 1-1
  - 1-Many
  - Many-Many
- **Foreign Keys?**



# One to One

- **Almost always embed 1-1 relationships in MongoDB**
  - User and UserProfile
  - User and AccountPreferences
  - User and Photo
- **Allows us to get all the info with one query**

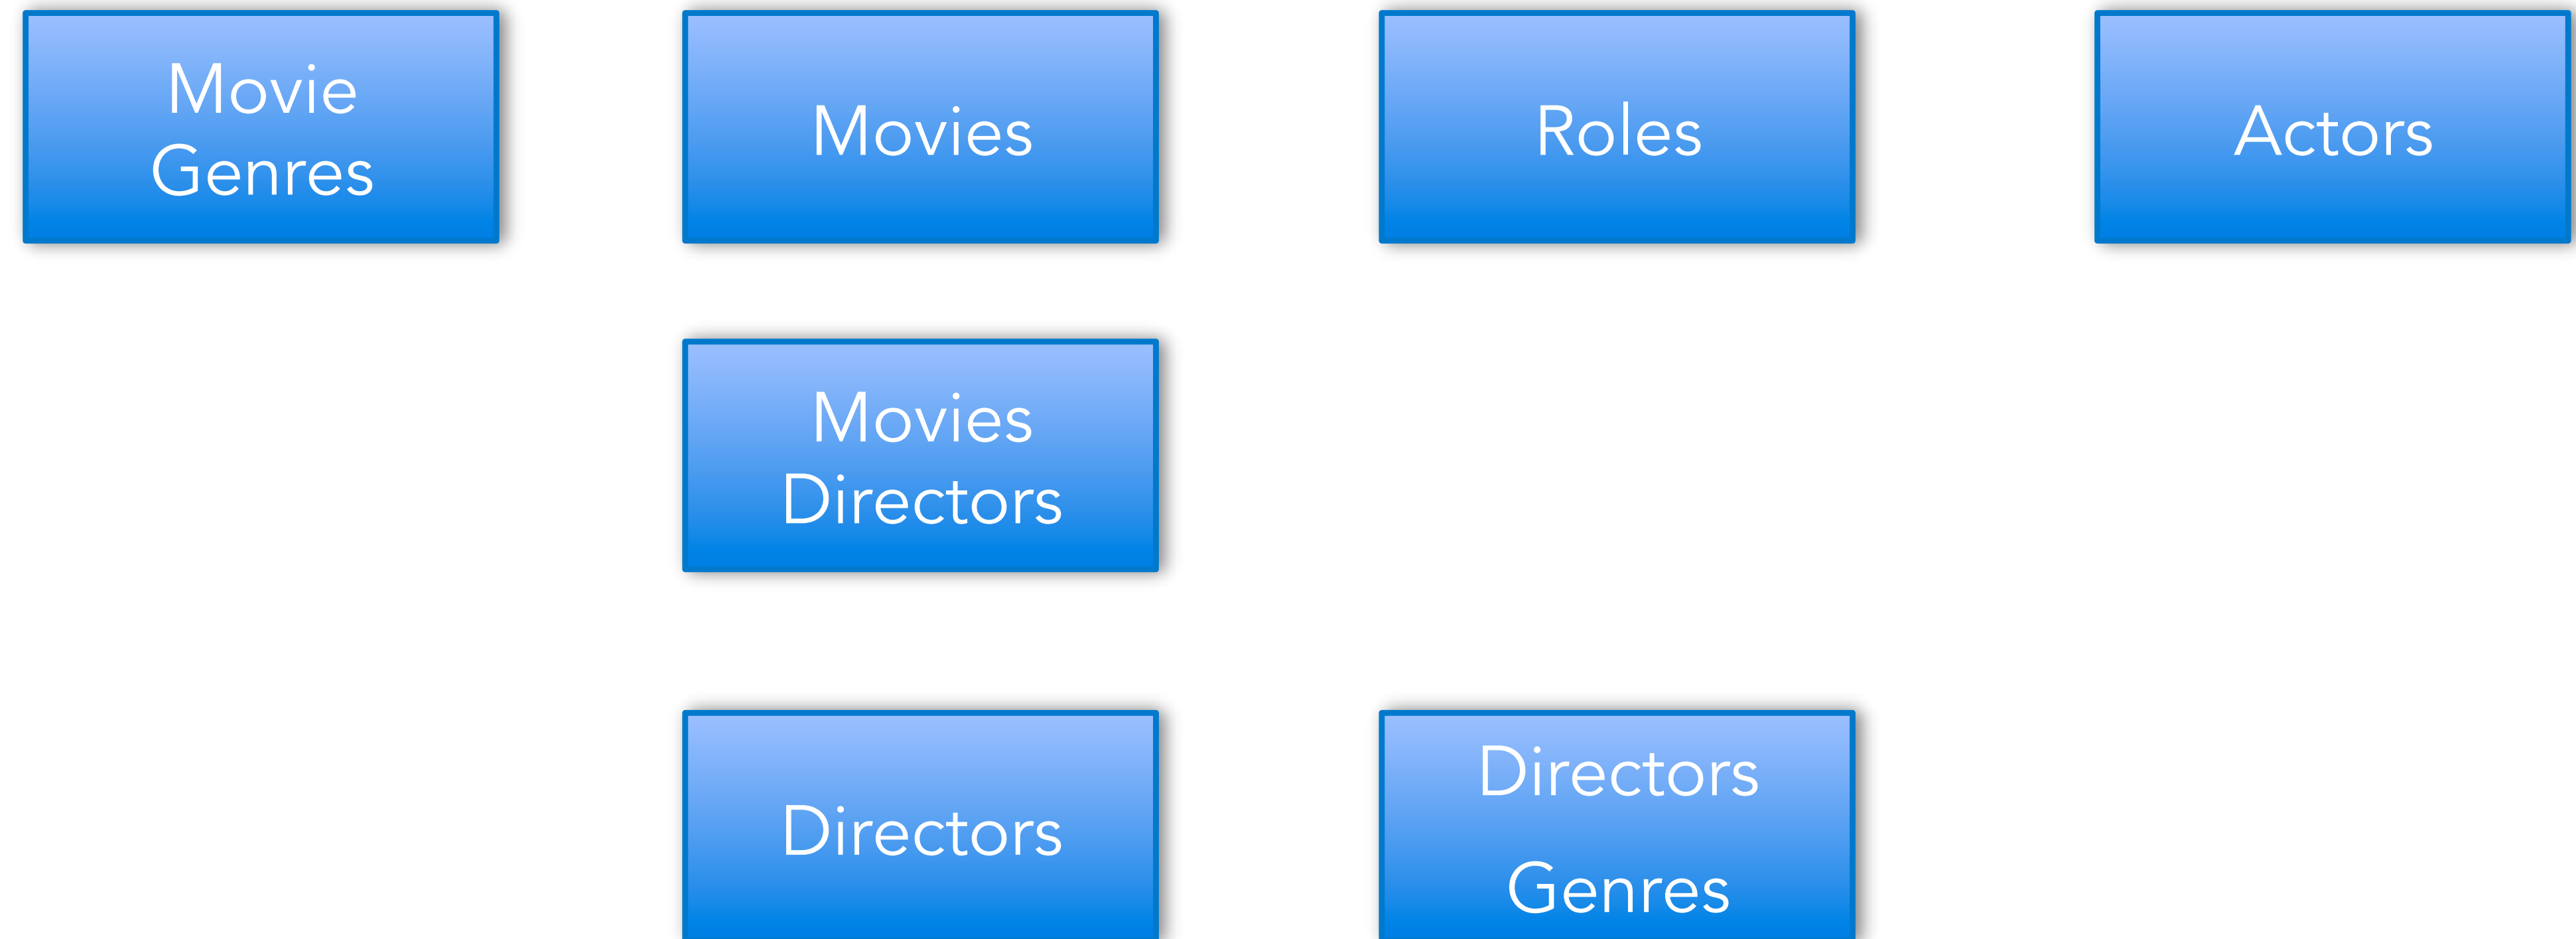
# One to Many

- Split into One-to-Few and One-to-Many
- One-to-Few
  - Prefer embedding
  - Examples: User and Addresses, Blog Post to Category
- One-to-Many
  - Prefer References
  - Examples: User and Posts, Blog Post to Comments
  - Who gets the reference? Typically the “many” (posts/comments in this example)

# Many to Many

- Social Networking's Many to Many problem
- Can embed followers or followees as references
- What happens on twitter when you have more than 16MB of references?
  - One document can reference the next document of followers
  - Linked list

# SQL to Mongo Conversion



# Movies to MovieGenres

- Few to Many
- From movies perspective, it's One to Few
- Recommendation: embed as string array

# Movies to Directors

- Many to Many
- From both movies and directors perspective, only a few
- Recommendation: embed in movies (why?)

# Movies to Actors

- Many to Many
- From both movies and directors perspective, only a few
- Where does the role information go?
- Recommendation: embed both in movies (why?)

# Workshop

## ● WikiStack

- Walk you through installing and using MongoDB
- Wikipedia for Fullstackers
- Application of everything we've learned so far
- We will do intermediate reviews to ensure everyone is on track